

Viola Zhong

Location: New York (can relocate) | **Email:** violazhongg@gmail.com | **Website** | [GitHub](#) | [Google Scholar](#)

ABOUT ME

From zero to first-author ICML workshop paper in 10 months, alongside 7 years of production engineering

PUBLICATIONS

[Refusal Lives Downstream of Persona in Chat Models, Mechanistic Interpretability Workshop, ICML, 2026](#)

- Showed persona gates refusal expression at late layers; refusal 97%→2% restored by persona knockout

[Framing Bias and Approaching Fairness in Algorithmic Tenant Screening Tools in NYCHRL, FAcCT Tutorial, 2019](#)

- Proposed a tenant score-challenge mechanism and legal-to-technical feature guidance

INDEPENDENT AI SAFETY RESEARCH

[AI-accelerated cyber exploitation: reframing the threat and a hybrid response](#) (April, 2026, Swift Centre)

- Proposed a shared vulnerability-discovery capability amortized across defenders, independently converging with core elements of Anthropic's Project Glasswing (announced the next day)

[I and Thou: Turning the Mirror on the Machine](#) (Feb, 2026)

- Built relational persona vectors that steer relational stance, found a cross-family safety-bypass vulnerability

[Llama2 Safety Evaluation Anatomy](#) (Jan, 2026)

- Mapped refusal in Llama-2-13b-chat as a two-stage detect-then-execute pipeline; projecting out the PC3 safety direction at layers 14-19 dropped refusal from 98.5% to 0%

[LLM Hallucinations: An Internal Tug of War](#) (Aug, 2025)

- Localized a unified certainty/uncertainty circuit overlapping the user-representation layers in Llama-2-13b-chat, independently converging with "incentivized guessing" in [OpenAI's hallucination paper](#)

TECHNICAL EXPERIENCE

Software Engineer, DoubleVerify, New York, NY (2019-)

- Architected an agentic product-ops ticket-resolution system, reducing human-in-the-loop handling by 70+%
- Designed and built an internal tool that reduced an incident impact analysis process from 3 hours to 5 mins

POLICY EXPERIENCE

Research Fellow, New York City Commission on Human Rights, New York, NY (2018-2019)

- Co-authored software developer best practices for building ML-driven applications with policymakers; published first-author tutorial in ACM FAcCT; self-directed research on fairness in machine learning
- Independently identified and analyzed legal vulnerabilities in CoreLogic's tenant-screening products; forecasted a lawsuit, [Arroyo v. CoreLogic \(2018\)](#), later confirmed predicted harms in June 2018

ENTREPRENEURIAL EXPERIENCE

Founding team, Smart Order, Hangzhou, China (2014-2016)

- Led marketing & partnership strategy; built 35 partnerships; grew subscribers 4x in one month

EDUCATION

New York University, New York, NY (2016-2017)

Master of Science in Data Science and Urban Informatics | GPA: 3.7/4.0

Zhejiang Gongshang University, Hangzhou, China (2010-2014)

Bachelor of Science in Economics | Mathematics GPA: 4.0/4.0

TECHNICAL SKILLS

- **Systems & ML stack:** Python, C#, Java, GCP, Kubernetes, mainstream DBs, PyTorch, vLLM, transformers
- **Interpretability methods:** linear probe, activation patching, activation steering, sparse autoencoders